

ICE: Exploratory Data Analysis (EDA) through Descriptive Statistics

Name: Cabot Steward

DATA 3300

In this assignment on exploratory data analysis (EDA) through descriptive statistics, we will start by examining the normal distribution using histogram plots and comparing it to non-normal distributions. Additionally, we'll cover descriptive statistics, including measures of central tendency, shape, and spread. Finally, we'll form hypotheses based on these descriptive statistics to compare different groups and draw meaningful insights from the data

Company Scenario

GridX, a Greentech firm, focuses on large-scale decarbonization by providing analytics on clean energy options to utility and energy companies. They create models for various transitions to clean energy and offer insights and advice on pricing. This helps companies understand the financial impact of adopting greener technologies. GridX's software suite includes programs for billing analytics, customer service, and bundling clean energy options. Currently GridX operates in the US, but is looking to expand to other countries.

Your task is to assist GridX in identifying which countries' utility and energy companies they should prioritize for marketing (beyond the US). Use descriptive statistics to explore the dataset, identify patterns, and develop hypotheses that can inform recommendations for GridX's sales team.

Dataset variables

- Country: Name of the country
- Region: Geographic region to which the country belongs

- SDGi (Sustainable Development Goals Index): The SDG Index evaluates countries based on indicators such as poverty, education, health, and environmental sustainability, providing a comprehensive overview of their efforts and achievements in sustainable development
- Life Expectancy: Average life expectancy in years
- HDI (Human Development Index): Human Development Index, a composite index measuring life expectancy, education, and per capita income
- Per Capita GDP: Gross Domestic Product (GDP) per capita.
- Income Group: Income group classification (e.g., LI for Low Income, UM for Upper Middle Income)
- Population: Country population
- Carbon Footprint: Ecological footprint of carbon emissions, measured in global hectares per person
- Total Ecological Footprint (Consumption): Total ecological footprint of consumption, measured in global hectares per person
- Ecological (Deficit) or Reserve: The difference between total ecological footprint and total biocapacity, indicating whether the country has an ecological deficit (negative value) or reserve (positive value)
- Number of Earths Required: The number of Earths that would be required if everyone lived like the average person in the country.
- Total Biocapacity: Total biocapacity, measured in global hectares per person

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
```

```
In [2]: df = pd.read_csv('Global Ecological Footprint 2023-1.csv', encoding='latin-1')
df.head()
```

Out[2]:

	Country	Region	SDGi	Life Exectancy	HDI	Per Capita GDP	Income Group	Population (millions)	Carbon Footprint	Total Ecological Footprint (Consumption)	Total biocapacity	E
0	Afghanistan	Middle East/Central Asia	52.5	62	0.48	NaN	LI	40.8	0.2	0.8	0.51	
1	Albania	Other Europe	71.6	76	0.8	14889	UM	2.9	0.7	2.1	1.18	
2	Algeria	Africa	71.5	76	0.75	11137	UM	45.4	1.2	2.2	0.66	
3	Angola	Africa	50.9	62	0.59	6304	LM	35.0	0.3	0.9	1.59	
4	Antigua and Barbuda	Central America/Caribbean	NaN	78	0.79	18749	HI	0.1	NaN	2.9	0.92	

1. If our main goal is to identify countries which should be prioritized by GridX, what should be some particular variables of interest for us to prioritize? Why?

- carbon footprint -> to know their current impact is significant
- gdp -> can they afford the solution
- total ecological footprint -
- population - per capita calculations
- income group
- country

2. Take a minute - come up with a question (or two) you'd be interested in investigating related to any of these variables of interest we've just listed above -- in the context of identifying countries that should be prioritized by GridX

1. Top 5 countries by carbon footprint per capita?
2. Correlation between income group and total ecological footprint?

We can begin to explore these types of questions by performing EDA through descriptive statistics. But before we know which descriptive stats to prioritize, it can be helpful to examine the distribution of a variable to know whether our observations follow a normal distribution (or not).

Begin by checking which variables can be evaluated

```
In [3]: # check info and vars
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 182 entries, 0 to 181
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               182 non-null    object
1   Region                                182 non-null    object
2   SDGi                                   159 non-null    object
3   Life Exectancy                         176 non-null    object
4   HDI                                     173 non-null    object
5   Per Capita GDP                         165 non-null    object
6   Income Group                           178 non-null    object
7   Population (millions)                  182 non-null    float64
8   Carbon Footprint                       152 non-null    float64
9   Total Ecological Footprint (Consumption) 181 non-null    float64
10  Total biocapacity                      181 non-null    float64
11  Ecological (Deficit) or Reserve          181 non-null    float64
12  Number of Earths required               181 non-null    float64
dtypes: float64(6), object(7)
memory usage: 18.6+ KB
```

Convert variables that should be floats!

```
In [4]: # convert as necessary
for col in ['SDGi', 'Life Exectancy', 'HDI', 'Per Capita GDP', ]:
    df[col] = pd.to_numeric(df[col], errors='coerce') # convert each column to a numeric if possible

# Print some info
df.info(), df.head()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 182 entries, 0 to 181
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Country	182 non-null	object
1	Region	182 non-null	object
2	SDGi	158 non-null	float64
3	Life Exectancy	175 non-null	float64
4	HDI	171 non-null	float64
5	Per Capita GDP	163 non-null	float64
6	Income Group	178 non-null	object
7	Population (millions)	182 non-null	float64
8	Carbon Footprint	152 non-null	float64
9	Total Ecological Footprint (Consumption)	181 non-null	float64
10	Total biocapacity	181 non-null	float64
11	Ecological (Deficit) or Reserve	181 non-null	float64
12	Number of Earths required	181 non-null	float64

```
dtypes: float64(10), object(3)
```

```
memory usage: 18.6+ KB
```

Out[4]: (None,

	Country	Region	SDGi	Life Exectancy	HDI	\
0	Afghanistan	Middle East/Central Asia	52.5	62.0	0.48	
1	Albania	Other Europe	71.6	76.0	0.80	
2	Algeria	Africa	71.5	76.0	0.75	
3	Angola	Africa	50.9	62.0	0.59	
4	Antigua and Barbuda	Central America/Caribbean	NaN	78.0	0.79	

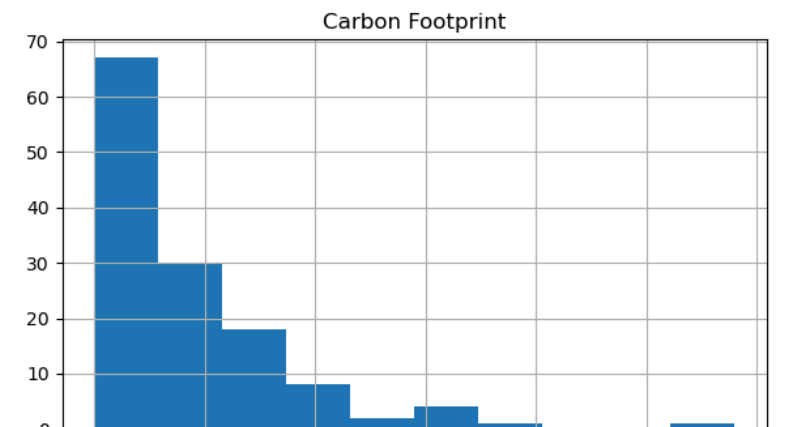
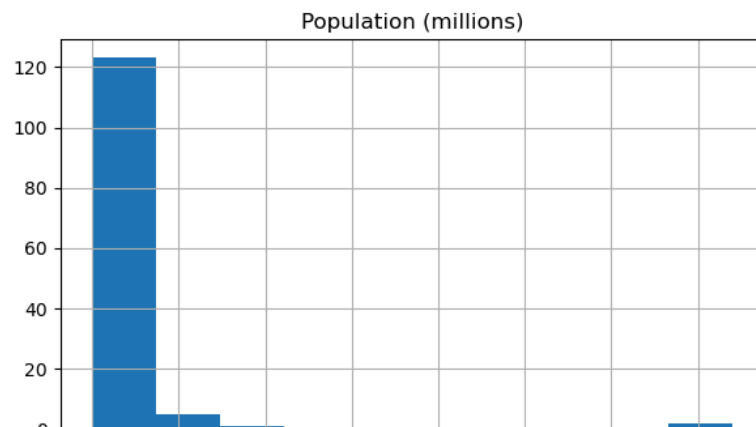
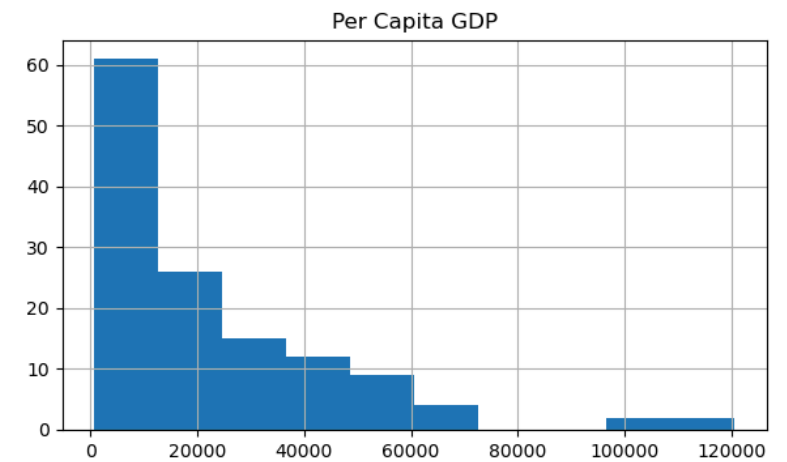
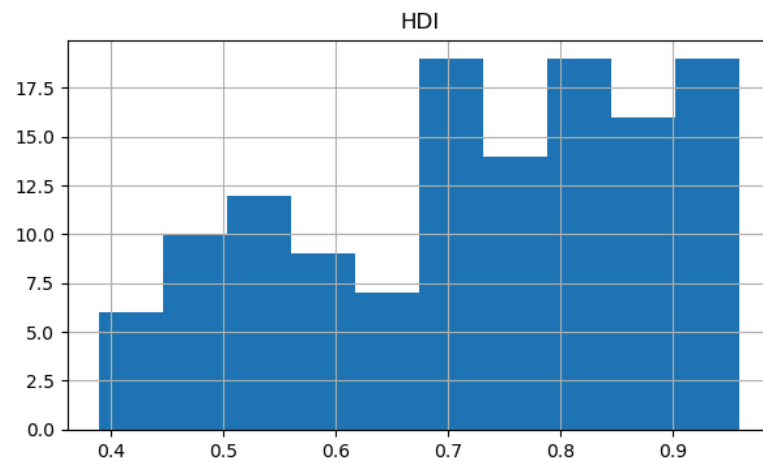
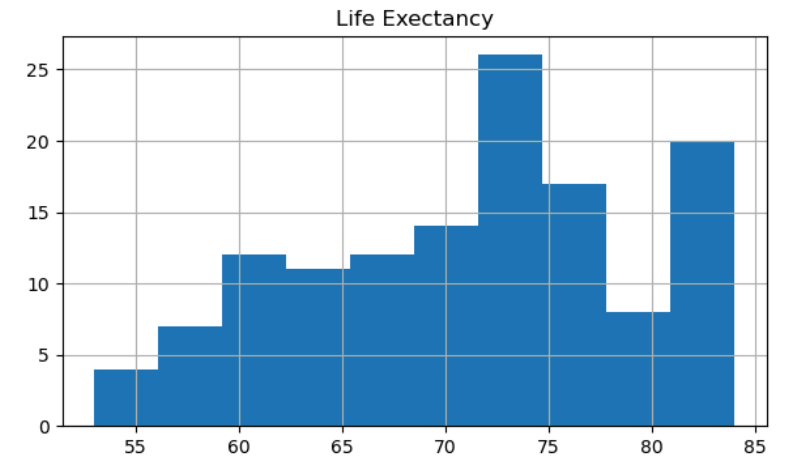

	Per Capita GDP	Income Group	Population (millions)	Carbon Footprint	\
0	NaN	LI	40.8	0.2	
1	14889.0	UM	2.9	0.7	
2	11137.0	UM	45.4	1.2	
3	6304.0	LM	35.0	0.3	
4	18749.0	HI	0.1	NaN	

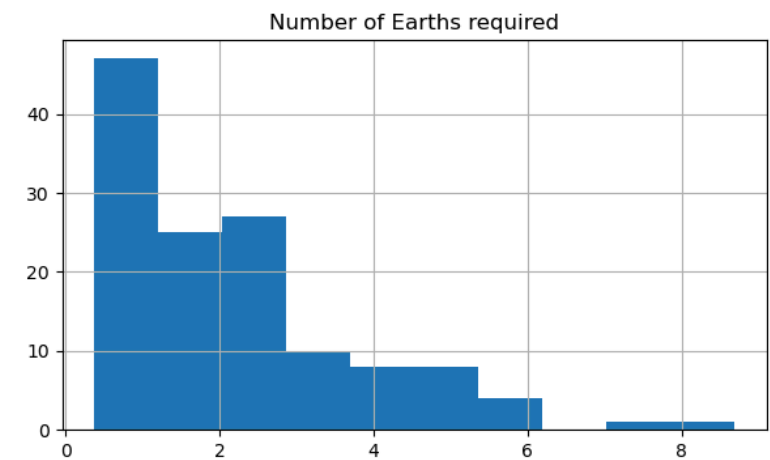
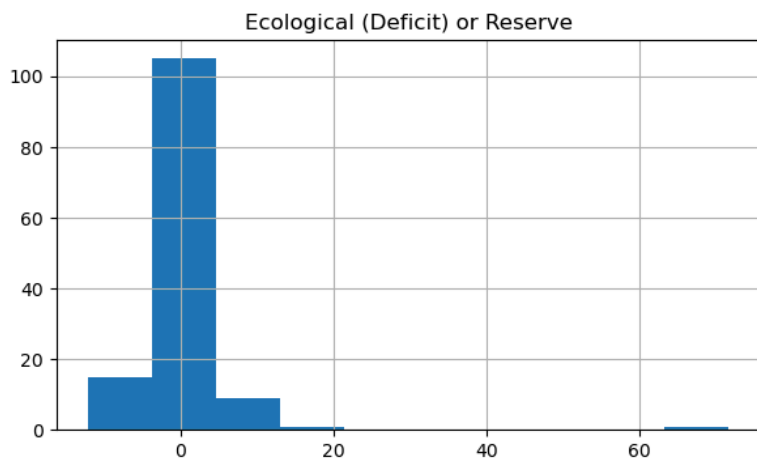
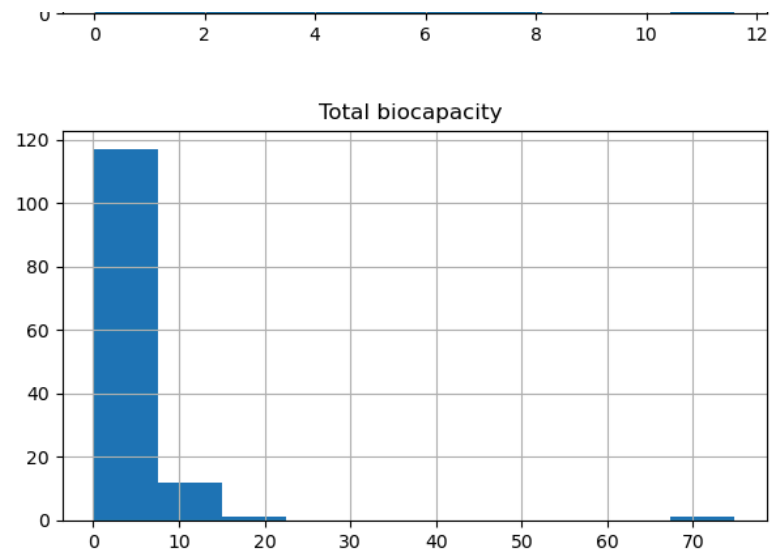
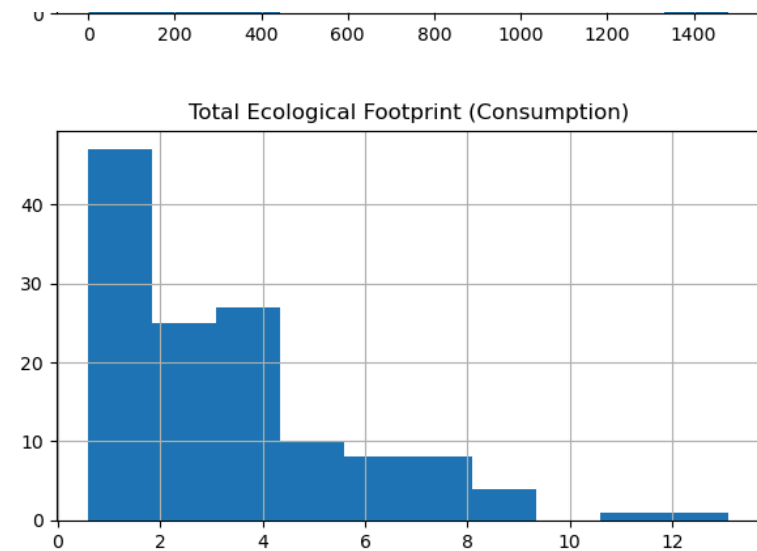
	Total Ecological Footprint (Consumption)	Total biocapacity	\
0	0.8	0.51	
1	2.1	1.18	
2	2.2	0.66	
3	0.9	1.59	
4	2.9	0.92	

	Ecological (Deficit) or Reserve	Number of Earths required
0	-0.29	0.53
1	-0.89	1.37
2	-1.56	1.47
3	0.73	0.57
4	-2.02	1.94

Visualize the Distributions

```
In [5]: df2 = df.dropna()
df2.hist(layout=(5, 2), figsize=(16, 24), bins=10) #generate histograms of remaining numerical vars
plt.show()
```





3. How would we describe the shape of the following distributions?

- SDGi: Left skewed, roughly unimodal
- HDI: not unimodal, it technically is quadmodal, but slightly left skewed if we had to.
- Carbon Footprint: right skewed, roughly unimodal
- Ecological Deficit: technically right skewed, leptokurtic

4. What implications do the shape of these distributions have for interpreting descriptive statistics? What measure of central tendency should NOT be relied on?

- Standard deviations will likely be skewed as there are outliers but we should utilize this! This is important
- On top of this average will be skewed so should be avoided if possible

Descriptive Statistics

Next, let's start examining some descriptive statistics. We'll begin by grouping our descriptive stats by Region, then we can also drill down to specific countries!

5. Return to your posed question(s) above. What information can you gather from this descriptive statistics table to address your question(s)?

- using the median values we can compare the regional differences between the carbon vs total ecological footprint and hypothesize that countries with carbon footprint.

```
In [6]: pd.set_option('display.max_columns', None) # display all columns in dataframe
df2.groupby('Region')['Carbon Footprint'].median().sort_values(ascending=False)
```

```
Out[6]: Region
North America          4.60
EU-27                  2.45
Middle East/Central Asia 2.40
Other Europe           1.80
Asia-Pacific           0.80
Central America/Caribbean 0.80
South America          0.70
Africa                 0.20
Name: Carbon Footprint, dtype: float64
```

6. Which region has the most spread in Carbon Footprint? What measure(s) indicate this?

```
In [7]: df2.groupby('Region')['Carbon Footprint'].std().sort_values(ascending=False)
```

```
Out[7]: Region
Middle East/Central Asia    2.972063
North America               2.080865
Asia-Pacific                1.856199
EU-27                      1.330202
Central America/Caribbean  1.212286
Other Europe                0.775743
South America               0.512522
Africa                      0.451457
Name: Carbon Footprint, dtype: float64
```

- The middle east/central asia has the most spread. We can see this from the standard deviation that is shown above.

7. Take a few minutes to explore the full descriptive stats table grouped by Region. Note any interesting pattern in your descriptive statistics table, and create a hypothesis as to the cause of this finding.

```
In [10]: df2[['Region', 'Population (millions)', 'Carbon Footprint']].groupby('Region').describe()
```

Out[10]:

Region	Population (millions)													
	count	mean	std	min	25%	50%	75%	max	count	mean	std	min	25%	
Africa	37.0	30.883784	41.040289	1.0	5.800	17.70	35.000	216.7	37.0	0.329730	0.451457	0.0	0.100	
Asia-Pacific	21.0	196.495238	422.222628	0.4	5.900	33.20	125.600	1480.6	21.0	1.738095	1.856199	0.1	0.300	
Central America/Caribbean	11.0	6.318182	5.606749	0.3	2.200	5.20	8.950	18.6	11.0	1.181818	1.212286	0.1	0.550	
EU-27	22.0	19.090909	23.745895	0.6	5.525	9.85	17.175	83.9	22.0	2.609091	1.330202	1.4	1.925	
Middle East/Central Asia	19.0	20.210526	25.850530	1.8	4.850	10.00	25.200	86.0	19.0	3.173684	2.972063	0.2	1.100	
North America	3.0	168.266667	151.563760	38.4	85.000	131.60	233.200	334.8	3.0	3.500000	2.080865	1.1	2.850	
Other Europe	10.0	25.080000	47.015785	0.6	2.975	4.80	9.225	145.8	10.0	1.920000	0.775743	0.7	1.575	
South America	8.0	48.225000	69.965333	0.6	10.825	26.50	47.375	215.4	8.0	0.962500	0.512522	0.5	0.575	
<div><div></div></div>														

Companies with a higher population don't always have a higher carbon footprint. EU-27 and Africa are prime examples

Let's Examine the Descriptive Stats of Two Regions we'd like to Compare...

8. What region would you recommend targeting and why? What descriptive stats are you basing this on and what variables?

```
In [13]: # compare two regions of interest
df2.groupby('Region')[['Population (millions)', 'Total Ecological Footprint (Consumption)']].describe().loc[['EU-27',
```

Out[13]:

	Population (millions)								Total Ecological Footprint (Consumption)							
	count	mean	std	min	25%	50%	75%	max	count	mean	std	min	25%	50%	75%	max
Region																
EU-27	22.0	19.090909	23.745895	0.6	5.525	9.85	17.175	83.9	22.0	5.263636	1.925653	2.7	3.925	4.55	6.2	11.0
Africa	37.0	30.883784	41.040289	1.0	5.800	17.70	35.000	216.7	37.0	1.367568	0.674559	0.6	0.900	1.10	1.7	3.2

lets narrow in on EU-27 because they have high variability (std) in their total ecological footprint, meaning it might be more clear which countries within this region should be targeted

Next, Let's drill down more into one specific region and identify the countries that we might target...

```
In [20]: df3 = df2[df2['Region'] == 'EU-27']
print(len(df3))
df3.head()
```

22

Out[20]:

	Country	Region	SDGi	Life Exectancy	HDI	Per Capita GDP	Income Group	Population (millions)	Carbon Footprint	Total Ecological Footprint (Consumption)	Total biocapacity	Ecological (Deficit) or Reserve
8	Austria	EU-27	82.3	81.0	0.92	55460.0	HI	9.1	3.2	5.6	2.89	-2.73
14	Belgium	EU-27	79.7	82.0	0.94	52749.0	HI	11.7	3.6	6.7	1.16	-5.53
24	Bulgaria	EU-27	74.3	72.0	0.80	24490.0	UM	6.8	1.9	3.6	3.08	-0.50
44	Czech Republic	EU-27	80.5	77.0	0.89	40707.0	HI	10.7	2.8	5.1	2.52	-2.56
45	Denmark	EU-27	85.6	81.0	0.95	59333.0	HI	5.8	2.9	7.3	4.22	-3.07



```

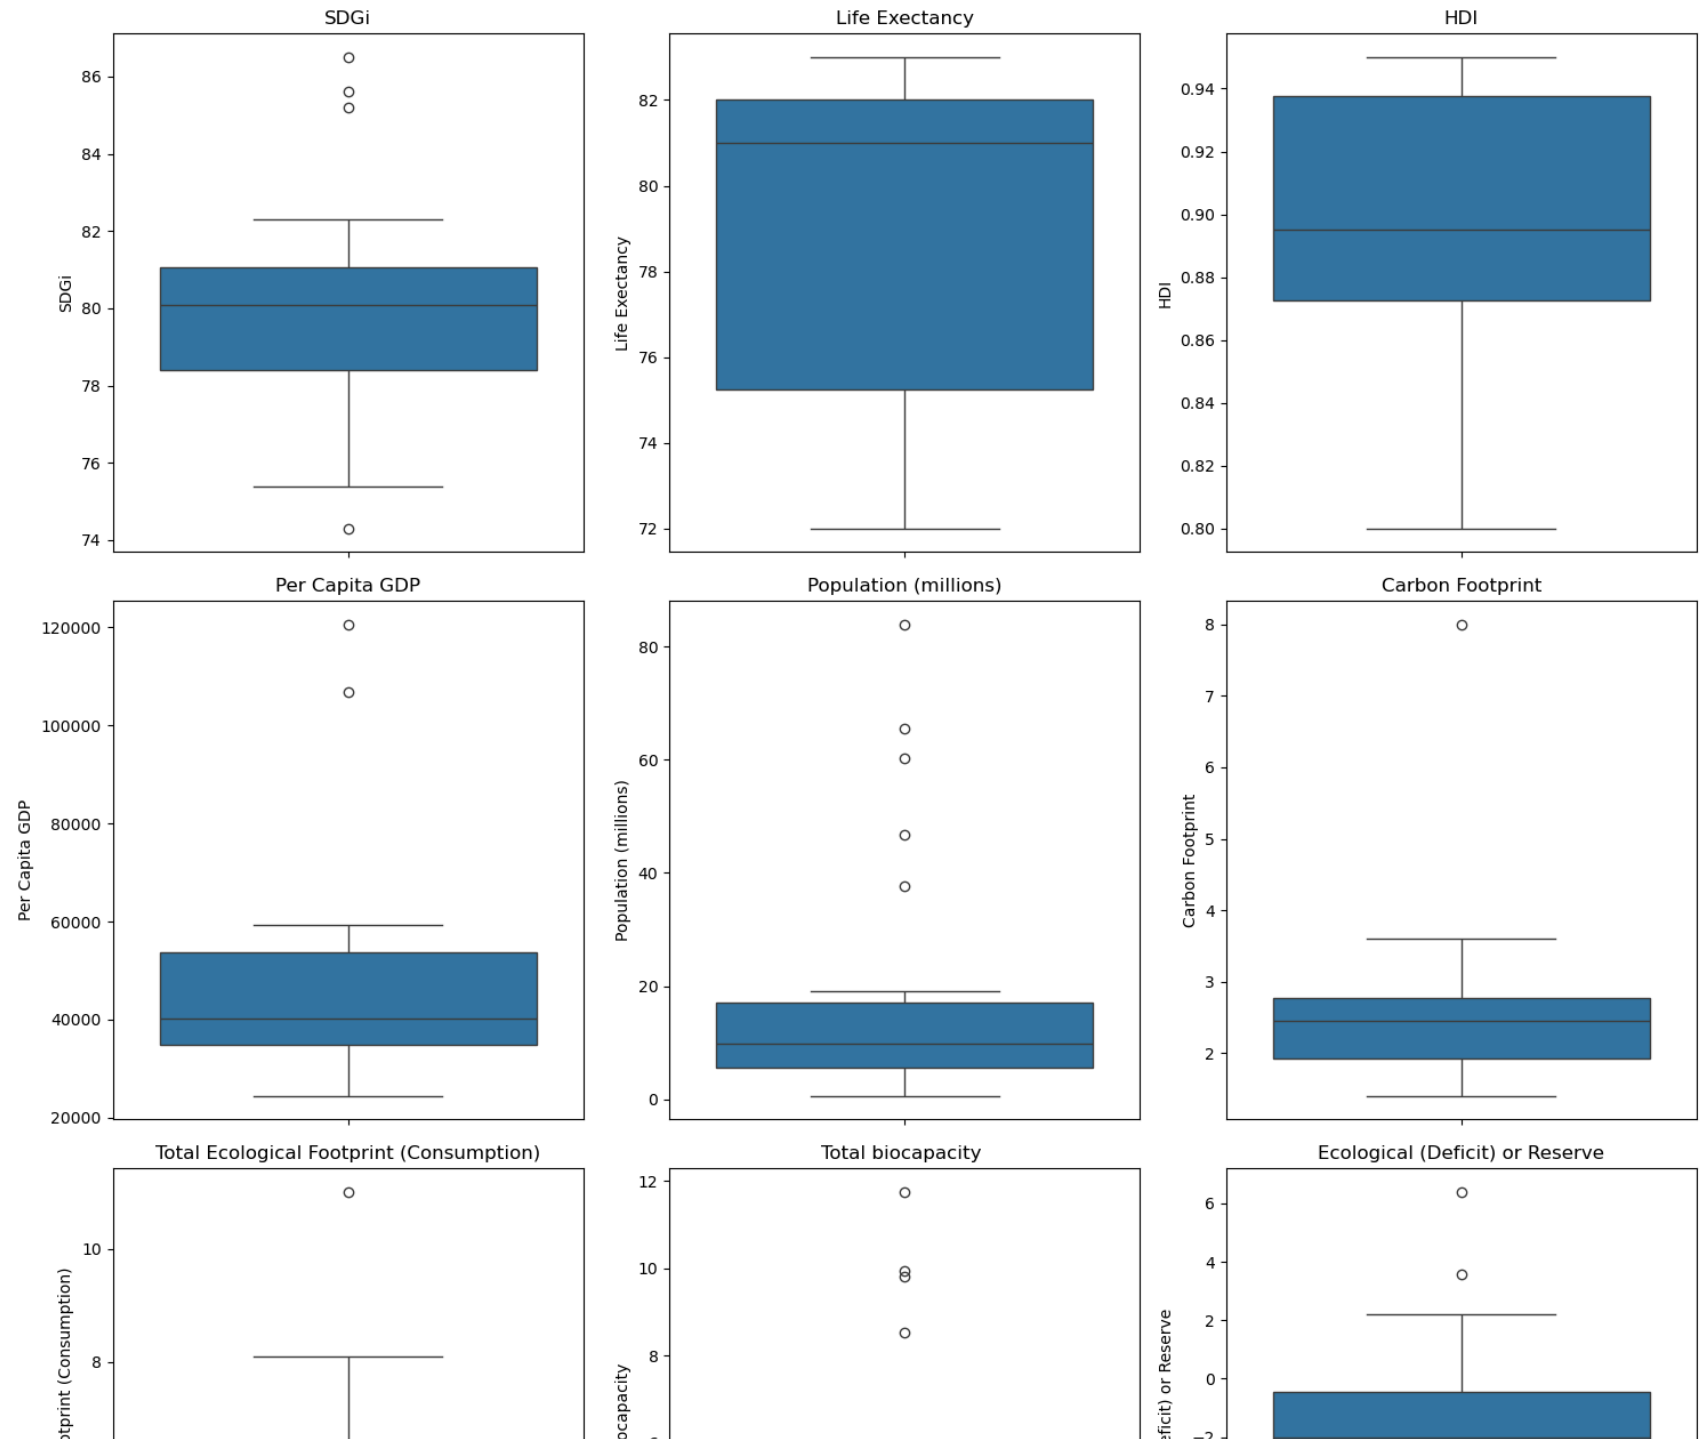
In [21]: numeric_cols = df3.select_dtypes(include=np.number).columns
num_cols = len(numeric_cols)
num_rows = (num_cols + 2) // 3 #

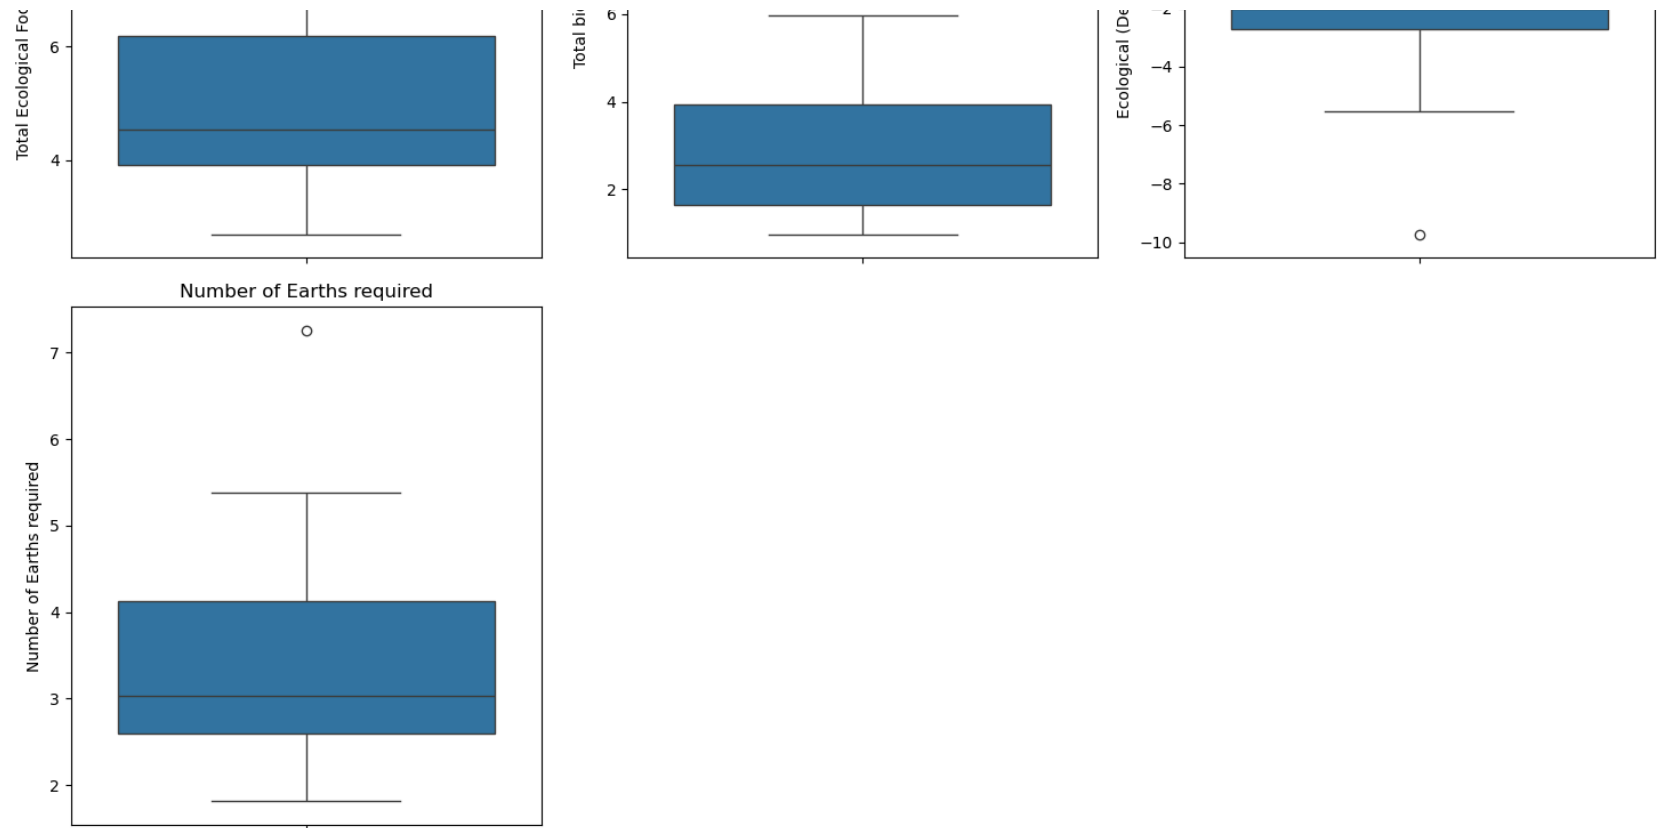
plt.figure(figsize=(15, 5 * num_rows)) # Add comments

for i, col in enumerate(numeric_cols):
    plt.subplot(num_rows, 3, i + 1) # Add comments
    sns.boxplot(y=df3[col])
    plt.title(col)

plt.tight_layout() # Add comments
plt.show()

```





9. What observations can we make from these distributions of just this region's countries that might help us narrow down on countries of interest?

gdp per capita

10. Visualize a relationship between two variables that investigates your hypothesis, originally posed questions, or that you'd like to explore based on observations above.

```
In [35]: plt.figure(figsize=(12, 10))

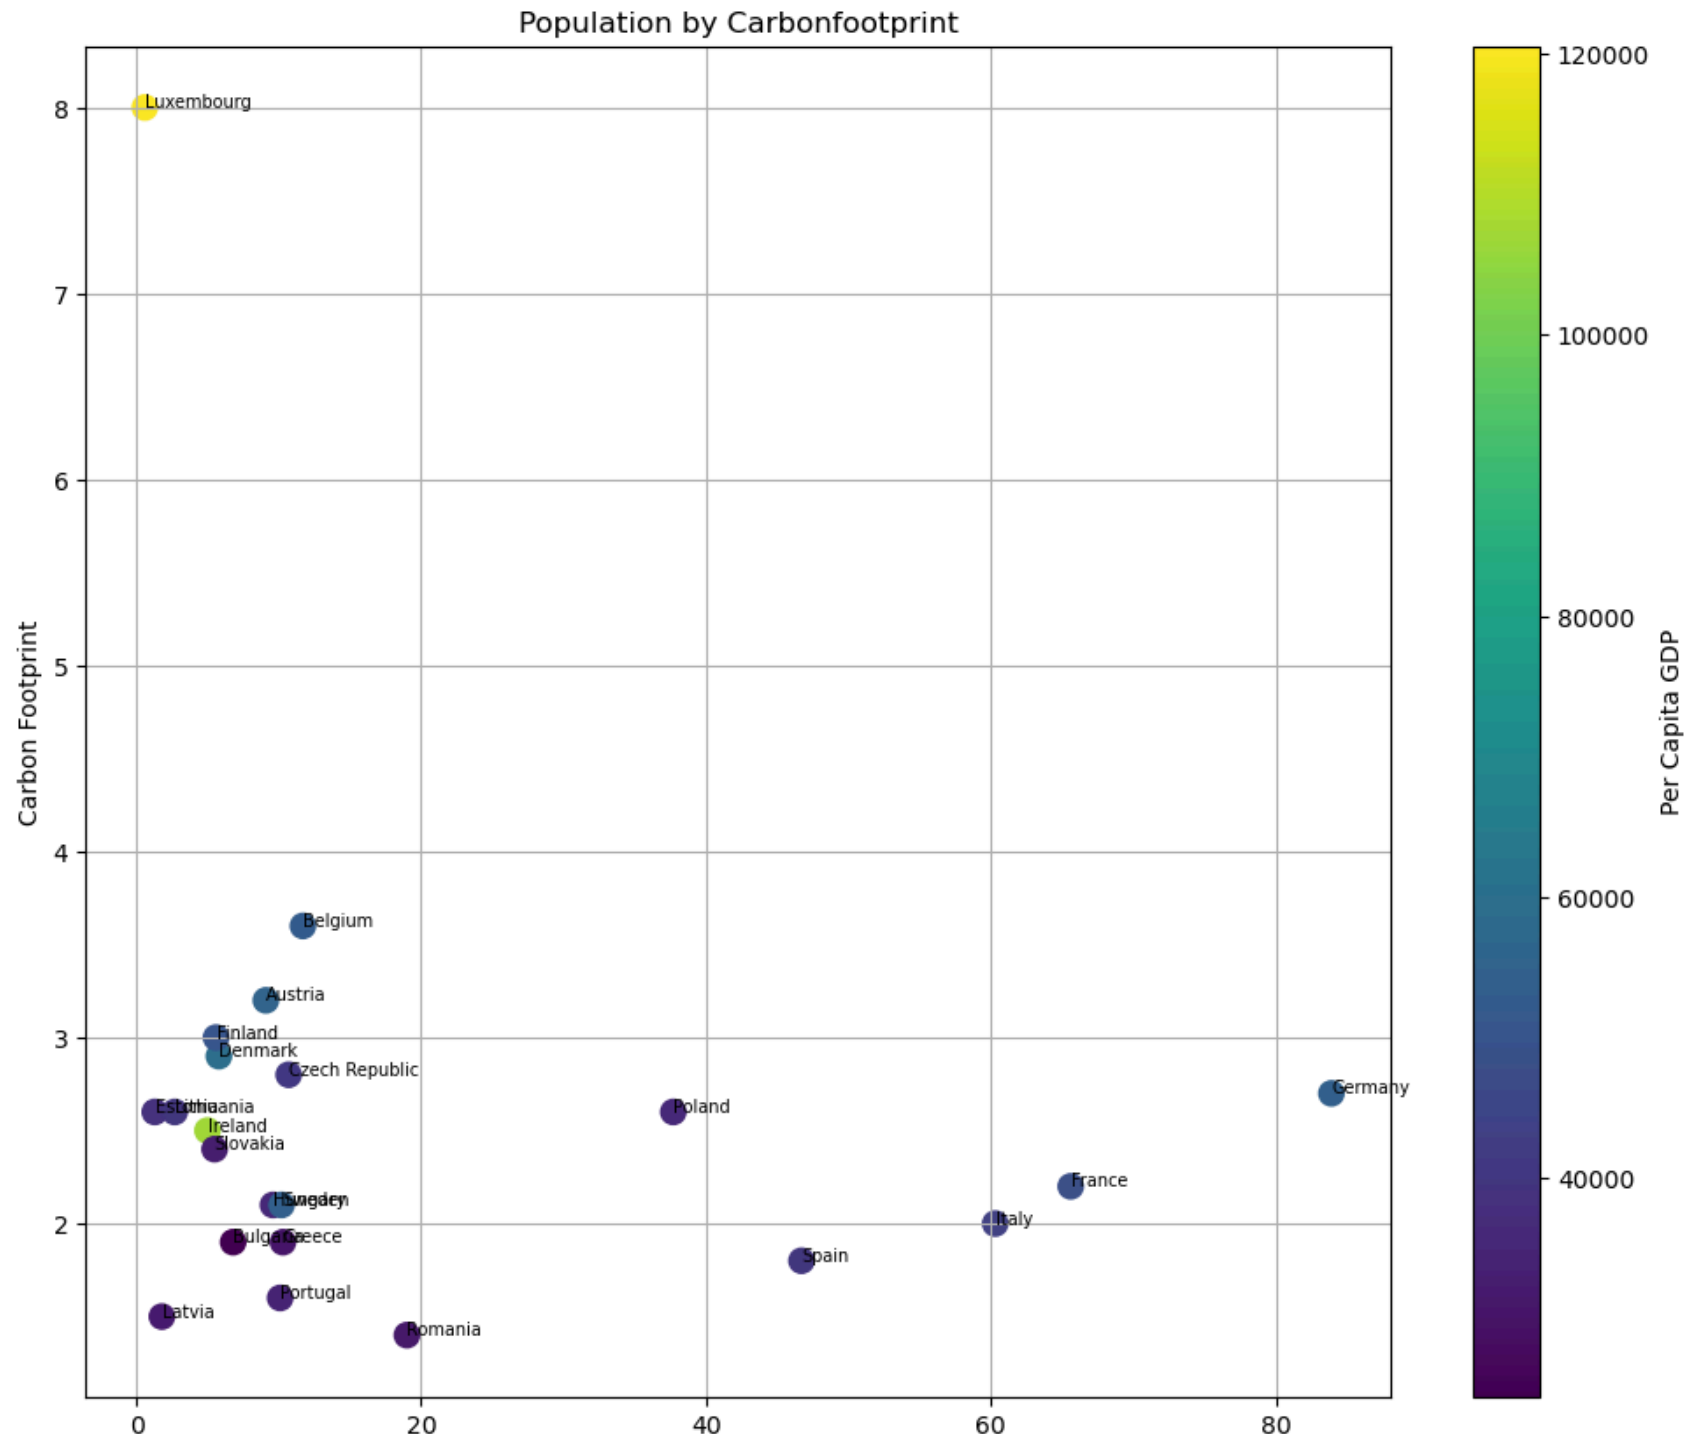
# Use a colormap for the scatter plot
scatter = plt.scatter(df3['Population (millions)'], df3['Carbon Footprint'],
                      c=df3['Per Capita GDP'], cmap='viridis', s=100)

# Add labels for each point
```

```
for i, row in df3.iterrows():
    plt.text(row['Population (millions)'], row['Carbon Footprint'], row['Country'], fontsize=7, ha='left')

# Add colorbar
colorbar = plt.colorbar(scatter)
colorbar.set_label('Per Capita GDP')

plt.title('Population by Carbonfootprint')
plt.xlabel('Population (millions)')
plt.ylabel('Carbon Footprint')
plt.grid(True)
plt.show()
```

Population (millions)**11. Based on your findings, which countries would you recommend GridX sales team focus on, and why?**

Luxembourg, Belgium, and Austria would be good ones to start the process working on. They have a medium size of population while having a decent amount of carbon footprint. You want to find a country that has a high gdp, and high carbon footprint. Germany is massive with a low carbon footprint.